IDMC MS4 Practice MCQ Question Bank

160 Questions | 100 Easy | 30 Medium | 30 Hard

Candidate Practice Edition - Based on the supplied MS4 IDMC MCQ source

Difficulty	Questions	Marks / Question	Total Marks
Easy	100	1	100
Medium	30	2	60
Hard	30	2	60
Total	160	-	220

Each question has one best answer. Explanations are provided for learning and review.

EASY QUESTIONS

100 questions | 1 mark each

Q1. Data Profiling

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A new customer feed has undocumented nulls and inconsistent formats. What should an IDMC developer do first to understand the data before writing cleansing logic?

A. Increase the Secure Agent memory before inspecting the data
B. Profile the source data before designing transformation rules
C. Create a taskflow schedule before checking the source
D. Define target load order without examining source values

Correct Answer: B. Profile the source data before designing transformation rules

Explanation: Profiling is used first to understand patterns, nulls, completeness, distributions, and other quality characteristics that can influence mapping design.

Q2. Data Profiling

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A mapping is failing because a source column contains unexpected value patterns. Which activity should be performed before changing the transformation rules?

A. Increase the Secure Agent memory before inspecting the data
B. Create a taskflow schedule before checking the source
C. Profile the source data before designing transformation rules
D. Define target load order without examining source values

Correct Answer: C. Profile the source data before designing transformation rules

Explanation: Profiling is used first to understand patterns, nulls, completeness, distributions, and other quality characteristics that can influence mapping design.

Q3. Data Profiling

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The team has received a new supplier dataset and wants evidence about completeness and value distributions before defining validation rules. What is the best first step?

A. Increase the Secure Agent memory before inspecting the data
B. Create a taskflow schedule before checking the source
C. Profile the source data before designing transformation rules
D. Define target load order without examining source values

Correct Answer: C. Profile the source data before designing transformation rules

Explanation: Profiling is used first to understand patterns, nulls, completeness, distributions, and other quality characteristics that can influence mapping design.

Q4. Data Profiling

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A developer wants to avoid designing transformation logic based only on assumptions about source quality. Which practice provides the most useful evidence?

A. Increase the Secure Agent memory before inspecting the data
B. Create a taskflow schedule before checking the source
C. Profile the source data before designing transformation rules
D. Define target load order without examining source values

Correct Answer: C. Profile the source data before designing transformation rules

Explanation: Profiling is used first to understand patterns, nulls, completeness, distributions, and other quality characteristics that can influence mapping design.

Q5. Expression Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A sales mapping must calculate NET_AMOUNT from quantity, unit price, and discount before loading the target. Which transformation best fits this requirement?

A. Use a Sorter transformation to derive calculated values
B. Use an Expression transformation to derive or modify output fields
C. Use a Router transformation to create arithmetic expressions
D. Use a Union transformation to modify individual field values

Correct Answer: B. Use an Expression transformation to derive or modify output fields

Explanation: The Expression transformation is the appropriate mapping component for field-level calculations, derivations, conditional expressions, and value manipulation.

Q6. Expression Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A customer pipeline needs to standardize a phone number and derive a status flag in the same mapping path. Which transformation should contain that field-level logic?

A. Use a Sorter transformation to derive calculated values
B. Use a Router transformation to create arithmetic expressions
C. Use a Union transformation to modify individual field values
D. Use an Expression transformation to derive or modify output fields

Correct Answer: D. Use an Expression transformation to derive or modify output fields

Explanation: The Expression transformation is the appropriate mapping component for field-level calculations, derivations, conditional expressions, and value manipulation.

Q7. Expression Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

An output column must be created from several incoming fields without grouping or joining rows. Which IDMC transformation should be selected?

A. Use a Sorter transformation to derive calculated values
B. Use an Expression transformation to derive or modify output fields
C. Use a Router transformation to create arithmetic expressions
D. Use a Union transformation to modify individual field values

Correct Answer: B. Use an Expression transformation to derive or modify output fields

Explanation: The Expression transformation is the appropriate mapping component for field-level calculations, derivations, conditional expressions, and value manipulation.

Q8. Expression Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The developer needs to convert an incoming timestamp to a reporting date and generate a derived business column. Which mapping component is most appropriate?

A. Use a Sorter transformation to derive calculated values
B. Use a Router transformation to create arithmetic expressions
C. Use an Expression transformation to derive or modify output fields
D. Use a Union transformation to modify individual field values

Correct Answer: C. Use an Expression transformation to derive or modify output fields

Explanation: The Expression transformation is the appropriate mapping component for field-level calculations, derivations, conditional expressions, and value manipulation.

Q9. Filter Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Only orders with STATUS = APPROVED should continue to the target. Which transformation should be placed in the mapping path?

A. Use a Filter transformation to keep only rows that satisfy the condition
B. Use an Aggregator transformation to remove rows by condition
C. Use a Joiner transformation to discard nonmatching business rows
D. Use a Union transformation to keep only qualifying records

Correct Answer: A. Use a Filter transformation to keep only rows that satisfy the condition

Explanation: A Filter transformation evaluates a row-level condition and passes only records that meet that condition.

Q10. Filter Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A healthcare feed contains active and inactive provider rows, but downstream processing needs active rows only. Which transformation should enforce this requirement?

A. Use an Aggregator transformation to remove rows by condition
B. Use a Joiner transformation to discard nonmatching business rows
C. Use a Union transformation to keep only qualifying records
D. Use a Filter transformation to keep only rows that satisfy the condition

Correct Answer: D. Use a Filter transformation to keep only rows that satisfy the condition

Explanation: A Filter transformation evaluates a row-level condition and passes only records that meet that condition.

Q11. Filter Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A mapping should stop low-value records from reaching later transformations based on a single row-level condition. Which component is best suited?

A. Use an Aggregator transformation to remove rows by condition
B. Use a Joiner transformation to discard nonmatching business rows
C. Use a Union transformation to keep only qualifying records
D. Use a Filter transformation to keep only rows that satisfy the condition

Correct Answer: D. Use a Filter transformation to keep only rows that satisfy the condition

Explanation: A Filter transformation evaluates a row-level condition and passes only records that meet that condition.

Q12. Filter Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The source delivers all transactions, while the target must receive only records from the current business date. Which transformation should perform this row selection?

A. Use an Aggregator transformation to remove rows by condition
B. Use a Joiner transformation to discard nonmatching business rows
C. Use a Union transformation to keep only qualifying records
D. Use a Filter transformation to keep only rows that satisfy the condition

Correct Answer: D. Use a Filter transformation to keep only rows that satisfy the condition

Explanation: A Filter transformation evaluates a row-level condition and passes only records that meet that condition.

Q13. Joiner Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Customer rows and support-ticket rows arrive through separate mapping pipelines and must be combined by CUSTOMER_ID. Which transformation should be used?

A. Use a Router transformation with multiple output groups
B. Use a Joiner transformation with the required matching condition
C. Use a Rank transformation on the matching columns
D. Use a Sequence Generator to align records

Correct Answer: B. Use a Joiner transformation with the required matching condition

Explanation: The Joiner transformation combines rows from two pipelines using configured join conditions.

Q14. Joiner Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Two incoming datasets do not come from the same source object but must be matched using a common business key. Which transformation is designed for this?

A. Use a Router transformation with multiple output groups
B. Use a Rank transformation on the matching columns
C. Use a Sequence Generator to align records
D. Use a Joiner transformation with the required matching condition

Correct Answer: D. Use a Joiner transformation with the required matching condition

Explanation: The Joiner transformation combines rows from two pipelines using configured join conditions.

Q15. Joiner Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A mapping requires fields from orders and shipment feeds to appear in one output row when ORDER_ID matches. Which transformation should perform the combination?

A. Use a Router transformation with multiple output groups
B. Use a Rank transformation on the matching columns
C. Use a Joiner transformation with the required matching condition
D. Use a Sequence Generator to align records

Correct Answer: C. Use a Joiner transformation with the required matching condition

Explanation: The Joiner transformation combines rows from two pipelines using configured join conditions.

Q16. Joiner Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The developer needs to combine two data streams based on an explicit equality condition rather than simply append their rows. Which component is appropriate?

A. Use a Router transformation with multiple output groups
B. Use a Rank transformation on the matching columns
C. Use a Sequence Generator to align records
D. Use a Joiner transformation with the required matching condition

Correct Answer: D. Use a Joiner transformation with the required matching condition

Explanation: The Joiner transformation combines rows from two pipelines using configured join conditions.

Q17. Aggregator Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A telecom mapping must calculate total call duration for each subscriber per day. Which transformation should perform the grouped calculation?

A. Use a Lookup transformation with grouped reference keys
B. Use a Router transformation with one group per total
C. Use an Aggregator transformation with the required group fields
D. Use a Sorter transformation to calculate grouped totals

Correct Answer: C. Use an Aggregator transformation with the required group fields

Explanation: Aggregator performs grouped calculations such as SUM, COUNT, MIN, MAX, and AVG using group-by fields.

Q18. Aggregator Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The target requires one row per branch containing COUNT and AVG values from many source transactions. Which transformation is the best fit?

A. Use an Aggregator transformation with the required group fields
B. Use a Lookup transformation with grouped reference keys
C. Use a Router transformation with one group per total
D. Use a Sorter transformation to calculate grouped totals

Correct Answer: A. Use an Aggregator transformation with the required group fields

Explanation: Aggregator performs grouped calculations such as SUM, COUNT, MIN, MAX, and AVG using group-by fields.

Q19. Aggregator Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A mapping must summarize claim amounts by policy type before loading an analytics table. Which transformation should be configured?

A. Use a Lookup transformation with grouped reference keys
B. Use a Router transformation with one group per total
C. Use an Aggregator transformation with the required group fields
D. Use a Sorter transformation to calculate grouped totals

Correct Answer: C. Use an Aggregator transformation with the required group fields

Explanation: Aggregator performs grouped calculations such as SUM, COUNT, MIN, MAX, and AVG using group-by fields.

Q20. Aggregator Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The business wants the maximum transaction amount for each customer rather than row-level detail. Which transformation supports this requirement directly?

A. Use a Lookup transformation with grouped reference keys
B. Use an Aggregator transformation with the required group fields
C. Use a Router transformation with one group per total
D. Use a Sorter transformation to calculate grouped totals

Correct Answer: B. Use an Aggregator transformation with the required group fields

Explanation: Aggregator performs grouped calculations such as SUM, COUNT, MIN, MAX, and AVG using group-by fields.

Q21. Router Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Transactions must be separated into Domestic, International, and Review paths inside one mapping. Which transformation is most suitable?

A. Use a Router transformation with separate output-group conditions
B. Use multiple Aggregator transformations for each category
C. Use a Union transformation with one condition per output
D. Use a Sorter transformation to create conditional branches

Correct Answer: A. Use a Router transformation with separate output-group conditions

Explanation: A Router evaluates conditions once and can send rows to multiple named output groups, which is appropriate for category-based branching.

Q22. Router Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A quality process needs Accepted, Rejected, and Manual Review outputs based on different conditions. Which transformation should implement the branching?

A. Use a Router transformation with separate output-group conditions
B. Use multiple Aggregator transformations for each category
C. Use a Union transformation with one condition per output
D. Use a Sorter transformation to create conditional branches

Correct Answer: A. Use a Router transformation with separate output-group conditions

Explanation: A Router evaluates conditions once and can send rows to multiple named output groups, which is appropriate for category-based branching.

Q23. Router Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

One incoming row may need to be evaluated against several business conditions and sent to the appropriate downstream group. Which transformation is designed for this?

A. Use multiple Aggregator transformations for each category
B. Use a Union transformation with one condition per output
C. Use a Sorter transformation to create conditional branches
D. Use a Router transformation with separate output-group conditions

Correct Answer: D. Use a Router transformation with separate output-group conditions

Explanation: A Router evaluates conditions once and can send rows to multiple named output groups, which is appropriate for category-based branching.

Q24. Router Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A mapping currently uses several duplicated filter branches for business categories. Which transformation can centralize the condition-based routing?

A. Use multiple Aggregator transformations for each category
B. Use a Union transformation with one condition per output
C. Use a Sorter transformation to create conditional branches
D. Use a Router transformation with separate output-group conditions

Correct Answer: D. Use a Router transformation with separate output-group conditions

Explanation: A Router evaluates conditions once and can send rows to multiple named output groups, which is appropriate for category-based branching.

Q25. Lookup Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Sales rows contain PRODUCT_ID but need CATEGORY and BRAND from a reference table. Which transformation should enrich the rows?

A. Use a Union transformation to append reference records
B. Use a Rank transformation to obtain reference attributes
C. Use a Lookup transformation to retrieve matching reference attributes
D. Use a Command task to query the reference object

Correct Answer: C. Use a Lookup transformation to retrieve matching reference attributes

Explanation: Lookup retrieves matching information from a reference object based on configured lookup conditions.

Q26. Lookup Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A mapping needs to translate a branch code into the corresponding region name using a reference object. Which transformation fits this use case?

A. Use a Lookup transformation to retrieve matching reference attributes
B. Use a Union transformation to append reference records
C. Use a Rank transformation to obtain reference attributes
D. Use a Command task to query the reference object

Correct Answer: A. Use a Lookup transformation to retrieve matching reference attributes

Explanation: Lookup retrieves matching information from a reference object based on configured lookup conditions.

Q27. Lookup Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Incoming policy records must be enriched with a risk class stored in a relatively small reference dataset. Which transformation should be selected?

A. Use a Union transformation to append reference records
B. Use a Rank transformation to obtain reference attributes
C. Use a Lookup transformation to retrieve matching reference attributes
D. Use a Command task to query the reference object

Correct Answer: C. Use a Lookup transformation to retrieve matching reference attributes

Explanation: Lookup retrieves matching information from a reference object based on configured lookup conditions.

Q28. Lookup Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The developer needs to retrieve one or more columns from a reference object based on a business key without combining two full pipelines. Which transformation is appropriate?

A. Use a Union transformation to append reference records
B. Use a Rank transformation to obtain reference attributes
C. Use a Lookup transformation to retrieve matching reference attributes
D. Use a Command task to query the reference object

Correct Answer: C. Use a Lookup transformation to retrieve matching reference attributes

Explanation: Lookup retrieves matching information from a reference object based on configured lookup conditions.

Q29. Sorter Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A downstream step expects invoice rows ordered by INVOICE_DATE descending. Which transformation should enforce the order?

A. Use a Rank transformation to sort the full stream only
B. Use an Aggregator transformation to control row order
C. Use a Lookup transformation to arrange records for the stated requirement
D. Use a Sorter transformation to order rows by the required fields

Correct Answer: D. Use a Sorter transformation to order rows by the required fields

Explanation: Sorter arranges records in ascending or descending order based on configured sort keys.

Q30. Sorter Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Before further processing, records must be arranged first by region and then by transaction date. Which transformation should be used?

A. Use a Rank transformation to sort the full stream only
B. Use an Aggregator transformation to control row order
C. Use a Sorter transformation to order rows by the required fields
D. Use a Lookup transformation to arrange records for the stated requirement

Correct Answer: C. Use a Sorter transformation to order rows by the required fields

Explanation: Sorter arranges records in ascending or descending order based on configured sort keys.

Q31. Sorter Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The mapping requirement is to control row order rather than filter, group, or join the data. Which transformation is intended for this?

A. Use a Rank transformation to sort the full stream only

B. Use a Sorter transformation to order rows by the required fields
C. Use an Aggregator transformation to control row order
D. Use a Lookup transformation to arrange records for the stated requirement

Correct Answer: B. Use a Sorter transformation to order rows by the required fields

Explanation: Sorter arranges records in ascending or descending order based on configured sort keys.

Q32. Sorter Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A file output must be produced in customer-number order for a legacy consumer. Which transformation should be added before the target?

A. Use a Rank transformation to sort the full stream only
B. Use an Aggregator transformation to control row order
C. Use a Sorter transformation to order rows by the required fields
D. Use a Lookup transformation to arrange records for the stated requirement

Correct Answer: C. Use a Sorter transformation to order rows by the required fields

Explanation: Sorter arranges records in ascending or descending order based on configured sort keys.

Q33. Union Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

North and South region feeds have the same field structure and must be combined into one stream before loading. Which transformation should be used?

A. Use a Joiner transformation without a join condition
B. Use a Router transformation to concatenate pipelines
C. Use a Union transformation to append compatible row streams
D. Use a Lookup transformation to stack compatible rows

Correct Answer: C. Use a Union transformation to append compatible row streams

Explanation: Union combines rows from multiple compatible pipelines into a single output stream without key-based matching.

Q34. Union Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Two compatible datasets must be appended row-by-row without matching on any business key. Which transformation is designed for this?

A. Use a Joiner transformation without a join condition
B. Use a Router transformation to concatenate pipelines
C. Use a Lookup transformation to stack compatible rows
D. Use a Union transformation to append compatible row streams

Correct Answer: D. Use a Union transformation to append compatible row streams

Explanation: Union combines rows from multiple compatible pipelines into a single output stream without key-based matching.

Q35. Union Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A mapping receives monthly files through separate paths and needs one consolidated row stream with the same columns. Which component is appropriate?

A. Use a Joiner transformation without a join condition
B. Use a Union transformation to append compatible row streams
C. Use a Router transformation to concatenate pipelines
D. Use a Lookup transformation to stack compatible rows

Correct Answer: B. Use a Union transformation to append compatible row streams

Explanation: Union combines rows from multiple compatible pipelines into a single output stream without key-based matching.

Q36. Union Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The requirement is to merge compatible pipelines vertically rather than combine columns from matching records. Which transformation should be selected?

A. Use a Joiner transformation without a join condition
B. Use a Router transformation to concatenate pipelines
C. Use a Union transformation to append compatible row streams
D. Use a Lookup transformation to stack compatible rows

Correct Answer: C. Use a Union transformation to append compatible row streams

Explanation: Union combines rows from multiple compatible pipelines into a single output stream without key-based matching.

Q37. Rank Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A sales report requires the five highest-value orders for each region. Which transformation should be considered?

A. Use an Aggregator transformation to return ranked row detail
B. Use a Rank transformation to select top or bottom records by ranking criteria
C. Use a Sorter transformation alone to restrict the output count
D. Use a Router transformation to calculate rank positions for the stated requirement

Correct Answer: B. Use a Rank transformation to select top or bottom records by ranking criteria

Explanation: Rank is designed to select top or bottom rows according to ranking criteria, optionally within groups.

Q38. Rank Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The business needs only the lowest three defect-rate records rather than the entire sorted dataset. Which transformation is the best fit?

A. Use an Aggregator transformation to return ranked row detail
B. Use a Rank transformation to select top or bottom records by ranking criteria
C. Use a Sorter transformation alone to restrict the output count
D. Use a Router transformation to calculate rank positions for the stated requirement

Correct Answer: B. Use a Rank transformation to select top or bottom records by ranking criteria

Explanation: Rank is designed to select top or bottom rows according to ranking criteria, optionally within groups.

Q39. Rank Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A mapping should return top-performing products according to revenue while preserving the selected row details. Which transformation supports this directly?

A. Use a Rank transformation to select top or bottom records by ranking criteria
B. Use an Aggregator transformation to return ranked row detail
C. Use a Sorter transformation alone to restrict the output count
D. Use a Router transformation to calculate rank positions for the stated requirement

Correct Answer: A. Use a Rank transformation to select top or bottom records by ranking criteria

Explanation: Rank is designed to select top or bottom rows according to ranking criteria, optionally within groups.

Q40. Rank Transformation

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The requirement is to choose a limited number of best or worst records based on a measure. Which transformation is intended for that pattern?

A. Use an Aggregator transformation to return ranked row detail
B. Use a Sorter transformation alone to restrict the output count
C. Use a Rank transformation to select top or bottom records by ranking criteria
D. Use a Router transformation to calculate rank positions for the stated requirement

Correct Answer: C. Use a Rank transformation to select top or bottom records by ranking criteria

Explanation: Rank is designed to select top or bottom rows according to ranking criteria, optionally within groups.

Q41. Mapplets and Reuse

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Five client mappings perform the same standardization and validation steps. What design reduces duplicated transformation logic?

A. Copy the same transformations into every mapping independently for the stated requirement
B. Hardcode the shared logic inside each mapping task within the integration workflow
C. Create a separate Secure Agent for each copy of the logic
D. Create reusable mapplet logic and invoke it from the mappings that need it

Correct Answer: D. Create reusable mapplet logic and invoke it from the mappings that need it

Explanation: Mapplets support reuse of a group of transformation logic across mappings, improving consistency and maintainability.

Q42. Mapplets and Reuse

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A shared cleansing rule changes frequently and must remain consistent across many mappings. Which reusable design should the team adopt?

A. Copy the same transformations into every mapping independently for the stated requirement
B. Create reusable mapplet logic and invoke it from the mappings that need it
C. Hardcode the shared logic inside each mapping task within the integration workflow
D. Create a separate Secure Agent for each copy of the logic

Correct Answer: B. Create reusable mapplet logic and invoke it from the mappings that need it

Explanation: Mapplets support reuse of a group of transformation logic across mappings, improving consistency and maintainability.

Q43. Mapplets and Reuse

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The development team wants one maintained component for a repeated group of transformations rather than many copied versions. What should they create?

A. Copy the same transformations into every mapping independently for the stated requirement
B. Create reusable mapplet logic and invoke it from the mappings that need it
C. Hardcode the shared logic inside each mapping task within the integration workflow
D. Create a separate Secure Agent for each copy of the logic

Correct Answer: B. Create reusable mapplet logic and invoke it from the mappings that need it

Explanation: Mapplets support reuse of a group of transformation logic across mappings, improving consistency and maintainability.

Q44. Mapplets and Reuse

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A reusable transformation sequence is required across several pipelines while allowing each mapping to keep its own source and target. Which design is most appropriate?

A. Create reusable mapplet logic and invoke it from the mappings that need it
B. Copy the same transformations into every mapping independently for the stated requirement
C. Hardcode the shared logic inside each mapping task within the integration workflow
D. Create a separate Secure Agent for each copy of the logic

Correct Answer: A. Create reusable mapplet logic and invoke it from the mappings that need it

Explanation: Mapplets support reuse of a group of transformation logic across mappings, improving consistency and maintainability.

Q45. Mapping and Mapping Task

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A pipeline design is complete, but the team still needs a runnable asset with runtime settings. Which design relationship is correct?

A. Put all transformation design only in the mapping task for the stated requirement
B. Use the mapping only for scheduling and the task for field logic
C. Keep transformation design in the mapping and use a mapping task for runnable execution settings
D. Replace the mapping task with a connection object for execution

Correct Answer: C. Keep transformation design in the mapping and use a mapping task for runnable execution settings

Explanation: A mapping defines source-to-target transformation logic, while a mapping task provides the runnable configuration such as runtime properties and parameter values.

Q46. Mapping and Mapping Task

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

An engineer is separating design-time logic from execution-time configuration. How should mapping and mapping task responsibilities be divided?

A. Put all transformation design only in the mapping task for the stated requirement
B. Use the mapping only for scheduling and the task for field logic
C. Keep transformation design in the mapping and use a mapping task for runnable execution settings
D. Replace the mapping task with a connection object for execution

Correct Answer: C. Keep transformation design in the mapping and use a mapping task for runnable execution settings

Explanation: A mapping defines source-to-target transformation logic, while a mapping task provides the runnable configuration such as runtime properties and parameter values.

Q47. Mapping and Mapping Task

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The source-to-target logic should remain reusable while schedules and runtime values can vary by execution. Which arrangement supports this?

A. Put all transformation design only in the mapping task for the stated requirement
B. Keep transformation design in the mapping and use a mapping task for runnable execution settings
C. Use the mapping only for scheduling and the task for field logic
D. Replace the mapping task with a connection object for execution

Correct Answer: B. Keep transformation design in the mapping and use a mapping task for runnable execution settings

Explanation: A mapping defines source-to-target transformation logic, while a mapping task provides the runnable configuration such as runtime properties and parameter values.

Q48. Mapping and Mapping Task

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A developer wants to change runtime properties without redesigning transformations. Which IDMC asset pairing enables that separation?

A. Put all transformation design only in the mapping task for the stated requirement
B. Keep transformation design in the mapping and use a mapping task for runnable execution settings
C. Use the mapping only for scheduling and the task for field logic
D. Replace the mapping task with a connection object for execution

Correct Answer: B. Keep transformation design in the mapping and use a mapping task for runnable execution settings

Explanation: A mapping defines source-to-target transformation logic, while a mapping task provides the runnable configuration such as runtime properties and parameter values.

Q49. Secure Agent

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A CDI mapping must read an on-premises database that is not directly exposed to the cloud. Which IDMC component normally provides the runtime connectivity?

A. Use a taskflow to replace runtime connectivity for the stated requirement
B. Use a mapplet as the execution runtime within the integration workflow
C. Use a dashboard to provide source-system connectivity as the primary design approach
D. Use the Secure Agent to execute jobs and connect IDMC to configured systems

Correct Answer: D. Use the Secure Agent to execute jobs and connect IDMC to configured systems

Explanation: The Secure Agent is the runtime component that executes integration jobs and provides connectivity to configured cloud or on-premises systems.

Q50. Secure Agent

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The design exists in IDMC, but the job still needs a runtime component to execute near the enterprise data source. Which component fulfills this role?

A. Use a taskflow to replace runtime connectivity for the stated requirement
B. Use a mapplet as the execution runtime within the integration workflow
C. Use a dashboard to provide source-system connectivity as the primary design approach
D. Use the Secure Agent to execute jobs and connect IDMC to configured systems

Correct Answer: D. Use the Secure Agent to execute jobs and connect IDMC to configured systems

Explanation: The Secure Agent is the runtime component that executes integration jobs and provides connectivity to configured cloud or on-premises systems.

Q51. Secure Agent

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

An organization needs IDMC jobs to access systems inside its network while retaining centrally managed cloud design. Which component is required?

A. Use a taskflow to replace runtime connectivity for the stated requirement
B. Use a mapplet as the execution runtime within the integration workflow
C. Use a dashboard to provide source-system connectivity as the primary design approach
D. Use the Secure Agent to execute jobs and connect IDMC to configured systems

Correct Answer: D. Use the Secure Agent to execute jobs and connect IDMC to configured systems

Explanation: The Secure Agent is the runtime component that executes integration jobs and provides connectivity to configured cloud or on-premises systems.

Q52. Secure Agent

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Which component bridges IDMC services with configured data systems during job execution rather than merely storing connection metadata?

A. Use a taskflow to replace runtime connectivity for the stated requirement
B. Use a mapplet as the execution runtime within the integration workflow

C. Use a dashboard to provide source-system connectivity as the primary design approach
D. Use the Secure Agent to execute jobs and connect IDMC to configured systems

Correct Answer: D. Use the Secure Agent to execute jobs and connect IDMC to configured systems

Explanation: The Secure Agent is the runtime component that executes integration jobs and provides connectivity to configured cloud or on-premises systems.

Q53. Runtime Environment

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A mapping task must be assigned to the correct group of Secure Agents for execution. Which IDMC object represents that execution choice?

A. Use a runtime environment that contains the Secure Agent resources for job execution
B. Use a project folder as the execution container for the stated requirement
C. Use a connection object as the execution engine within the integration workflow
D. Use a mapplet to group Secure Agents as the primary design approach

Correct Answer: A. Use a runtime environment that contains the Secure Agent resources for job execution

Explanation: A runtime environment identifies the runtime resources, including one or more Secure Agents, that execute IDMC jobs.

Q54. Runtime Environment

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The platform team wants jobs to use a defined set of runtime resources rather than an arbitrary agent. Which configuration should they select?

A. Use a runtime environment that contains the Secure Agent resources for job execution
B. Use a project folder as the execution container for the stated requirement
C. Use a connection object as the execution engine within the integration workflow
D. Use a mapplet to group Secure Agents as the primary design approach

Correct Answer: A. Use a runtime environment that contains the Secure Agent resources for job execution

Explanation: A runtime environment identifies the runtime resources, including one or more Secure Agents, that execute IDMC jobs.

Q55. Runtime Environment

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A task is ready to run, but the administrator must specify where it will execute. Which setting identifies the appropriate Secure Agent resources?

A. Use a project folder as the execution container for the stated requirement
B. Use a connection object as the execution engine within the integration workflow
C. Use a runtime environment that contains the Secure Agent resources for job execution
D. Use a mapplet to group Secure Agents as the primary design approach

Correct Answer: C. Use a runtime environment that contains the Secure Agent resources for job execution

Explanation: A runtime environment identifies the runtime resources, including one or more Secure Agents, that execute IDMC jobs.

Q56. Runtime Environment

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Several Secure Agents are managed together for integration workloads. Which IDMC construct is used to associate them with job execution?

A. Use a runtime environment that contains the Secure Agent resources for job execution
B. Use a project folder as the execution container for the stated requirement
C. Use a connection object as the execution engine within the integration workflow
D. Use a mapplet to group Secure Agents as the primary design approach

Correct Answer: A. Use a runtime environment that contains the Secure Agent resources for job execution

Explanation: A runtime environment identifies the runtime resources, including one or more Secure Agents, that execute IDMC jobs.

Q57. Connections

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A new Oracle source must be added to CDI. Where should host, service, user, and authentication details be configured?

A. Configure a connection with the endpoint and authentication properties for the system
B. Store the endpoint only in a taskflow decision step
C. Use a mapplet as the source-system credential store
D. Use a monitoring dashboard to define database access

Correct Answer: A. Configure a connection with the endpoint and authentication properties for the system

Explanation: Connections hold the endpoint and authentication properties used by IDMC to access source and target systems.

Q58. Connections

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The same target system will be used by several mappings. Which object should centrally define how IDMC accesses that system?

A. Configure a connection with the endpoint and authentication properties for the system
B. Store the endpoint only in a taskflow decision step
C. Use a mapplet as the source-system credential store
D. Use a monitoring dashboard to define database access

Correct Answer: A. Configure a connection with the endpoint and authentication properties for the system

Explanation: Connections hold the endpoint and authentication properties used by IDMC to access source and target systems.

Q59. Connections

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A developer needs a reusable definition of endpoint and login properties for a cloud database. Which IDMC asset should be created?

A. Configure a connection with the endpoint and authentication properties for the system
B. Store the endpoint only in a taskflow decision step
C. Use a mapplet as the source-system credential store
D. Use a monitoring dashboard to define database access

Correct Answer: A. Configure a connection with the endpoint and authentication properties for the system

Explanation: Connections hold the endpoint and authentication properties used by IDMC to access source and target systems.

Q60. Connections

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Before a source object can be selected in a mapping, IDMC needs access details for that source platform. Which object supplies them?

A. Configure a connection with the endpoint and authentication properties for the system
B. Store the endpoint only in a taskflow decision step
C. Use a mapplet as the source-system credential store
D. Use a monitoring dashboard to define database access

Correct Answer: A. Configure a connection with the endpoint and authentication properties for the system

Explanation: Connections hold the endpoint and authentication properties used by IDMC to access source and target systems.

Q61. Parameterization

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Development, test, and production use different file paths, but the transformation logic is identical. What design avoids three hardcoded mappings?

A. Use parameters so runtime values can change without rewriting the mapping logic
B. Duplicate the mapping for every runtime value for the stated requirement
C. Hardcode environment-specific values inside transformations within the integration workflow
D. Create a new Secure Agent whenever a value changes

Correct Answer: A. Use parameters so runtime values can change without rewriting the mapping logic

Explanation: Parameters make mappings and tasks flexible by allowing runtime values such as paths, table names, and connections to vary without redesign.

Q62. Parameterization

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A single mapping should load a different target table depending on the execution context. Which design technique provides this flexibility?

A. Duplicate the mapping for every runtime value for the stated requirement

B. Use parameters so runtime values can change without rewriting the mapping logic
C. Hardcode environment-specific values inside transformations within the integration workflow
D. Create a new Secure Agent whenever a value changes

Correct Answer: B. Use parameters so runtime values can change without rewriting the mapping logic

Explanation: Parameters make mappings and tasks flexible by allowing runtime values such as paths, table names, and connections to vary without redesign.

Q63. Parameterization

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The business wants to change a source file name at runtime without editing and republishing the mapping. Which feature should be used?

A. Use parameters so runtime values can change without rewriting the mapping logic
B. Duplicate the mapping for every runtime value for the stated requirement
C. Hardcode environment-specific values inside transformations within the integration workflow
D. Create a new Secure Agent whenever a value changes

Correct Answer: A. Use parameters so runtime values can change without rewriting the mapping logic

Explanation: Parameters make mappings and tasks flexible by allowing runtime values such as paths, table names, and connections to vary without redesign.

Q64. Parameterization

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A reusable pipeline must accept environment-specific values while keeping one maintained design. What is the appropriate IDMC approach?

A. Duplicate the mapping for every runtime value for the stated requirement
B. Hardcode environment-specific values inside transformations within the integration workflow
C. Use parameters so runtime values can change without rewriting the mapping logic
D. Create a new Secure Agent whenever a value changes

Correct Answer: C. Use parameters so runtime values can change without rewriting the mapping logic

Explanation: Parameters make mappings and tasks flexible by allowing runtime values such as paths, table names, and connections to vary without redesign.

Q65. Sequential Taskflow

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A pipeline must ingest data, validate it, and then load it, with each stage requiring the previous stage to finish. Which orchestration pattern is appropriate?

A. Use parallel execution even when every step depends on prior output
B. Use a Router transformation to orchestrate tasks for the stated requirement
C. Use a Union transformation to control task order within the integration workflow
D. Use sequential taskflow execution when each step depends on completion of the previous step

Correct Answer: D. Use sequential taskflow execution when each step depends on completion of the previous step

Explanation: Sequential taskflow design is appropriate when downstream tasks depend on upstream completion or output.

Q66. Sequential Taskflow

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Three mappings have strict dependencies and cannot overlap in execution. Which taskflow behavior should be configured?

A. Use parallel execution even when every step depends on prior output
B. Use sequential taskflow execution when each step depends on completion of the previous step
C. Use a Router transformation to orchestrate tasks for the stated requirement
D. Use a Union transformation to control task order within the integration workflow

Correct Answer: B. Use sequential taskflow execution when each step depends on completion of the previous step

Explanation: Sequential taskflow design is appropriate when downstream tasks depend on upstream completion or output.

Q67. Sequential Taskflow

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A downstream task consumes a file produced by the prior task. Which taskflow execution style best preserves the dependency?

A. Use sequential taskflow execution when each step depends on completion of the previous step
B. Use parallel execution even when every step depends on prior output
C. Use a Router transformation to orchestrate tasks for the stated requirement
D. Use a Union transformation to control task order within the integration workflow

Correct Answer: A. Use sequential taskflow execution when each step depends on completion of the previous step

Explanation: Sequential taskflow design is appropriate when downstream tasks depend on upstream completion or output.

Q68. Sequential Taskflow

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The workflow must guarantee a fixed order because each step requires results from the preceding step. Which orchestration approach should be used?

A. Use parallel execution even when every step depends on prior output
B. Use sequential taskflow execution when each step depends on completion of the previous step
C. Use a Router transformation to orchestrate tasks for the stated requirement
D. Use a Union transformation to control task order within the integration workflow

Correct Answer: B. Use sequential taskflow execution when each step depends on completion of the previous step

Explanation: Sequential taskflow design is appropriate when downstream tasks depend on upstream completion or output.

Q69. Parallel Taskflow

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Four regional mappings have no dependencies between them and the batch window is tight. Which taskflow design can reduce total elapsed time?

A. Force all independent tasks to execute one after another
B. Use a Lookup transformation to run tasks concurrently
C. Use a mapping task parameter to simulate parallelism
D. Use parallel execution paths for tasks that are independent of one another

Correct Answer: D. Use parallel execution paths for tasks that are independent of one another

Explanation: Parallel paths reduce elapsed time when tasks have no dependency on one another and can safely execute at the same time.

Q70. Parallel Taskflow

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Two independent validation tasks currently wait for each other even though neither uses the other's output. Which orchestration change is appropriate?

A. Use parallel execution paths for tasks that are independent of one another
B. Force all independent tasks to execute one after another
C. Use a Lookup transformation to run tasks concurrently
D. Use a mapping task parameter to simulate parallelism

Correct Answer: A. Use parallel execution paths for tasks that are independent of one another

Explanation: Parallel paths reduce elapsed time when tasks have no dependency on one another and can safely execute at the same time.

Q71. Parallel Taskflow

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A workflow contains several independent branches that can safely execute at the same time. Which taskflow capability should be used?

A. Force all independent tasks to execute one after another
B. Use a Lookup transformation to run tasks concurrently
C. Use a mapping task parameter to simulate parallelism
D. Use parallel execution paths for tasks that are independent of one another

Correct Answer: D. Use parallel execution paths for tasks that are independent of one another

Explanation: Parallel paths reduce elapsed time when tasks have no dependency on one another and can safely execute at the same time.

Q72. Parallel Taskflow

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The team wants to shorten an orchestration run without changing any mapping logic, and the tasks are independent. Which execution pattern should be chosen?

A. Force all independent tasks to execute one after another
B. Use a Lookup transformation to run tasks concurrently

C. Use parallel execution paths for tasks that are independent of one another
D. Use a mapping task parameter to simulate parallelism

Correct Answer: C. Use parallel execution paths for tasks that are independent of one another

Explanation: Parallel paths reduce elapsed time when tasks have no dependency on one another and can safely execute at the same time.

Q73. Decision Step

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

After validation, the workflow must load data only when the validation status is PASS; otherwise it should send an alert. Which taskflow feature is appropriate?

A. Use a Sorter transformation to choose the next task for the stated requirement
B. Use a decision or conditional step to select the execution path from runtime results
C. Use a Union transformation for conditional orchestration within the integration workflow
D. Use only a fixed sequential path regardless of outcomes as the primary design approach

Correct Answer: B. Use a decision or conditional step to select the execution path from runtime results

Explanation: A decision step evaluates a condition and directs taskflow execution to the appropriate branch.

Q74. Decision Step

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A taskflow needs one path for zero-row input and another path for nonzero input. Which orchestration component should control the branch?

A. Use a Sorter transformation to choose the next task for the stated requirement
B. Use a Union transformation for conditional orchestration within the integration workflow
C. Use a decision or conditional step to select the execution path from runtime results
D. Use only a fixed sequential path regardless of outcomes as the primary design approach

Correct Answer: C. Use a decision or conditional step to select the execution path from runtime results

Explanation: A decision step evaluates a condition and directs taskflow execution to the appropriate branch.

Q75. Decision Step

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The next task depends on a value returned by an earlier step. Which taskflow feature should evaluate that value and choose the path?

A. Use a Sorter transformation to choose the next task for the stated requirement
B. Use a decision or conditional step to select the execution path from runtime results
C. Use a Union transformation for conditional orchestration within the integration workflow
D. Use only a fixed sequential path regardless of outcomes as the primary design approach

Correct Answer: B. Use a decision or conditional step to select the execution path from runtime results

Explanation: A decision step evaluates a condition and directs taskflow execution to the appropriate branch.

Q76. Decision Step

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A workflow must behave differently for success, warning, and exception conditions. Which taskflow design concept enables conditional branching?

A. Use a Sorter transformation to choose the next task for the stated requirement
B. Use a decision or conditional step to select the execution path from runtime results
C. Use a Union transformation for conditional orchestration within the integration workflow
D. Use only a fixed sequential path regardless of outcomes as the primary design approach

Correct Answer: B. Use a decision or conditional step to select the execution path from runtime results

Explanation: A decision step evaluates a condition and directs taskflow execution to the appropriate branch.

Q77. Fault Handling

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

An external service sometimes fails temporarily, and the workflow needs a controlled response rather than an abrupt stop. Which taskflow capability should be designed?

A. Configure fault handling so failures are captured and handled without uncontrolled workflow termination
B. Ignore every failure and continue without recording it for the stated requirement
C. Duplicate every mapping to handle errors manually within the integration workflow
D. Use a Sorter transformation as the failure-recovery mechanism as the primary design approach

Correct Answer: A. Configure fault handling so failures are captured and handled without uncontrolled workflow termination

Explanation: Fault handling provides controlled recovery behavior such as retries, alternate paths, notifications, or other defined responses to errors.

Q78. Fault Handling

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A production orchestration must record failures and execute a defined recovery path. Which feature addresses this requirement?

A. Ignore every failure and continue without recording it for the stated requirement
B. Configure fault handling so failures are captured and handled without uncontrolled workflow termination
C. Duplicate every mapping to handle errors manually within the integration workflow
D. Use a Sorter transformation as the failure-recovery mechanism as the primary design approach

Correct Answer: B. Configure fault handling so failures are captured and handled without uncontrolled workflow termination

Explanation: Fault handling provides controlled recovery behavior such as retries, alternate paths, notifications, or other defined responses to errors.

Q79. Fault Handling

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The team wants predictable behavior when a step fails instead of leaving recovery to manual intervention. What should be configured?

A. Ignore every failure and continue without recording it for the stated requirement

B. Duplicate every mapping to handle errors manually within the integration workflow

C. Use a Sorter transformation as the failure-recovery mechanism as the primary design approach

D. Configure fault handling so failures are captured and handled without uncontrolled workflow termination

Correct Answer: D. Configure fault handling so failures are captured and handled without uncontrolled workflow termination

Explanation: Fault handling provides controlled recovery behavior such as retries, alternate paths, notifications, or other defined responses to errors.

Q80. Fault Handling

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A taskflow needs to distinguish successful execution from a failure condition and perform recovery actions. Which capability supports this?

A. Ignore every failure and continue without recording it for the stated requirement

B. Duplicate every mapping to handle errors manually within the integration workflow

C. Use a Sorter transformation as the failure-recovery mechanism as the primary design approach

D. Configure fault handling so failures are captured and handled without uncontrolled workflow termination

Correct Answer: D. Configure fault handling so failures are captured and handled without uncontrolled workflow termination

Explanation: Fault handling provides controlled recovery behavior such as retries, alternate paths, notifications, or other defined responses to errors.

Q81. Command and Event Steps

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A taskflow must invoke a shell-level utility as part of orchestration. Which category of taskflow capability is relevant?

A. Use an Aggregator transformation to execute operating-system commands for the stated requirement

B. Use a Joiner transformation to wait for external events within the integration workflow

C. Use the appropriate taskflow command or event capability for system-level actions or event-driven execution

D. Use a Lookup transformation to start system commands as the primary design approach

Correct Answer: C. Use the appropriate taskflow command or event capability for system-level actions or event-driven execution

Explanation: Taskflow command steps support system-level commands, while triggers/listeners support event-driven execution patterns.

Q82. Command and Event Steps

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The workflow should begin when a defined event signal arrives instead of only at a fixed time. Which orchestration capability should be considered?

A. Use the appropriate taskflow command or event capability for system-level actions or event-driven execution

B. Use an Aggregator transformation to execute operating-system commands for the stated requirement

C. Use a Joiner transformation to wait for external events within the integration workflow

D. Use a Lookup transformation to start system commands as the primary design approach

Correct Answer: A. Use the appropriate taskflow command or event capability for system-level actions or event-driven execution

Explanation: Taskflow command steps support system-level commands, while triggers/listeners support event-driven execution patterns.

Q83. Command and Event Steps

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A process needs to react to an incoming event and then start downstream integration steps. Which type of taskflow design supports this behavior?

A. Use an Aggregator transformation to execute operating-system commands for the stated requirement

B. Use the appropriate taskflow command or event capability for system-level actions or event-driven execution

C. Use a Joiner transformation to wait for external events within the integration workflow

D. Use a Lookup transformation to start system commands as the primary design approach

Correct Answer: B. Use the appropriate taskflow command or event capability for system-level actions or event-driven execution

Explanation: Taskflow command steps support system-level commands, while triggers/listeners support event-driven execution patterns.

Q84. Command and Event Steps

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The orchestration includes a system command outside the mapping itself. Where should this action be modeled?

A. Use an Aggregator transformation to execute operating-system commands for the stated requirement

B. Use the appropriate taskflow command or event capability for system-level actions or event-driven execution

C. Use a Joiner transformation to wait for external events within the integration workflow

D. Use a Lookup transformation to start system commands as the primary design approach

Correct Answer: B. Use the appropriate taskflow command or event capability for system-level actions or event-driven execution

Explanation: Taskflow command steps support system-level commands, while triggers/listeners support event-driven execution patterns.

Q85. Change Data Capture

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A source table changes continuously, but each integration run currently reloads millions of unchanged rows. Which processing approach should replace the full reload?

A. Reload the full source on every cycle regardless of changes
B. Use change data capture or incremental processing so only new and changed records are handled
C. Use result caching as the change-detection mechanism for the stated requirement
D. Use a Sorter transformation to identify changed rows within the integration workflow

Correct Answer: B. Use change data capture or incremental processing so only new and changed records are handled

Explanation: CDC/incremental processing reduces unnecessary work by focusing processing on inserts, updates, and deletes since the previous cycle.

Q86. Change Data Capture

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The target needs frequent updates while processing only records modified since the last successful run. Which pattern should be used?

A. Use change data capture or incremental processing so only new and changed records are handled
B. Reload the full source on every cycle regardless of changes
C. Use result caching as the change-detection mechanism for the stated requirement
D. Use a Sorter transformation to identify changed rows within the integration workflow

Correct Answer: A. Use change data capture or incremental processing so only new and changed records are handled

Explanation: CDC/incremental processing reduces unnecessary work by focusing processing on inserts, updates, and deletes since the previous cycle.

Q87. Change Data Capture

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

An order pipeline needs to capture inserts, updates, and deletes without repeatedly reading the entire source dataset. Which design is appropriate?

A. Reload the full source on every cycle regardless of changes
B. Use change data capture or incremental processing so only new and changed records are handled
C. Use result caching as the change-detection mechanism for the stated requirement
D. Use a Sorter transformation to identify changed rows within the integration workflow

Correct Answer: B. Use change data capture or incremental processing so only new and changed records are handled

Explanation: CDC/incremental processing reduces unnecessary work by focusing processing on inserts, updates, and deletes since the previous cycle.

Q88. Change Data Capture

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Data volume is growing, but only a small percentage changes between runs. Which integration strategy best reduces redundant processing?

A. Reload the full source on every cycle regardless of changes
B. Use result caching as the change-detection mechanism for the stated requirement
C. Use a Sorter transformation to identify changed rows within the integration workflow
D. Use change data capture or incremental processing so only new and changed records are handled

Correct Answer: D. Use change data capture or incremental processing so only new and changed records are handled

Explanation: CDC/incremental processing reduces unnecessary work by focusing processing on inserts, updates, and deletes since the previous cycle.

Q89. SQL Pushdown

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A mapping moves a very large dataset to the integration engine before filtering most of it out. Which optimization should be considered?

A. Move more data into the mapping engine before applying transformations
B. Disable database-side execution for joins and aggregations for the stated requirement
C. Increase data movement so the Secure Agent performs all operations
D. Push compatible transformation work to the database so processing occurs close to the data

Correct Answer: D. Push compatible transformation work to the database so processing occurs close to the data

Explanation: Pushdown improves performance by executing compatible operations in the database, reducing data movement and engine-side processing.

Q90. SQL Pushdown

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Snowflake compute is available, but joins and aggregations are still consuming heavy Secure Agent processing. Which execution strategy can reduce data movement?

A. Move more data into the mapping engine before applying transformations
B. Disable database-side execution for joins and aggregations for the stated requirement
C. Push compatible transformation work to the database so processing occurs close to the data
D. Increase data movement so the Secure Agent performs all operations

Correct Answer: C. Push compatible transformation work to the database so processing occurs close to the data

Explanation: Pushdown improves performance by executing compatible operations in the database, reducing data movement and engine-side processing.

Q91. SQL Pushdown

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The team wants filtering and aggregation to execute where the data already resides instead of in the mapping engine. Which optimization concept applies?

A. Move more data into the mapping engine before applying transformations
B. Disable database-side execution for joins and aggregations for the stated requirement

C. Increase data movement so the Secure Agent performs all operations

D. Push compatible transformation work to the database so processing occurs close to the data

Correct Answer: D. Push compatible transformation work to the database so processing occurs close to the data

Explanation: Pushdown improves performance by executing compatible operations in the database, reducing data movement and engine-side processing.

Q92. SQL Pushdown

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A pipeline should make better use of database processing power while preserving the same business logic. Which IDMC performance feature is relevant?

A. Push compatible transformation work to the database so processing occurs close to the data

B. Move more data into the mapping engine before applying transformations

C. Disable database-side execution for joins and aggregations for the stated requirement

D. Increase data movement so the Secure Agent performs all operations

Correct Answer: A. Push compatible transformation work to the database so processing occurs close to the data

Explanation: Pushdown improves performance by executing compatible operations in the database, reducing data movement and engine-side processing.

Q93. Batch and Real-Time Ingestion

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A nightly file drop can be processed at a fixed time, while fraud events must be handled continuously. Which ingestion design correctly distinguishes the two?

A. Use batch only for continuously arriving low-latency events for the stated requirement

B. Use real-time ingestion only for monthly scheduled files within the integration workflow

C. Choose batch for scheduled grouped loads and real-time ingestion for continuously arriving data that needs low latency

D. Treat batch and real-time as identical execution patterns as the primary design approach

Correct Answer: C. Choose batch for scheduled grouped loads and real-time ingestion for continuously arriving data that needs low latency

Explanation: Batch ingestion processes data in scheduled groups, while real-time ingestion is appropriate for continuously arriving data requiring low-latency handling.

Q94. Batch and Real-Time Ingestion

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The business accepts a once-per-day load for one feed but requires immediate processing for another. Which principle should guide the ingestion choice?

A. Choose batch for scheduled grouped loads and real-time ingestion for continuously arriving data that needs low latency

B. Use batch only for continuously arriving low-latency events for the stated requirement

C. Use real-time ingestion only for monthly scheduled files within the integration workflow

D. Treat batch and real-time as identical execution patterns as the primary design approach

Correct Answer: A. Choose batch for scheduled grouped loads and real-time ingestion for continuously arriving data that needs low latency

Explanation: Batch ingestion processes data in scheduled groups, while real-time ingestion is appropriate for continuously arriving data requiring low-latency handling.

Q95. Batch and Real-Time Ingestion

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A team is deciding between periodic grouped processing and continuous event processing. Which statement correctly describes the appropriate use of each?

A. Use batch only for continuously arriving low-latency events for the stated requirement
B. Use real-time ingestion only for monthly scheduled files within the integration workflow
C. Treat batch and real-time as identical execution patterns as the primary design approach
D. Choose batch for scheduled grouped loads and real-time ingestion for continuously arriving data that needs low latency

Correct Answer: D. Choose batch for scheduled grouped loads and real-time ingestion for continuously arriving data that needs low latency

Explanation: Batch ingestion processes data in scheduled groups, while real-time ingestion is appropriate for continuously arriving data requiring low-latency handling.

Q96. Batch and Real-Time Ingestion

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

Monthly reference files and continuously arriving transaction events have different latency requirements. Which ingestion approach should be assigned to each workload?

A. Use batch only for continuously arriving low-latency events for the stated requirement
B. Use real-time ingestion only for monthly scheduled files within the integration workflow
C. Treat batch and real-time as identical execution patterns as the primary design approach
D. Choose batch for scheduled grouped loads and real-time ingestion for continuously arriving data that needs low latency

Correct Answer: D. Choose batch for scheduled grouped loads and real-time ingestion for continuously arriving data that needs low latency

Explanation: Batch ingestion processes data in scheduled groups, while real-time ingestion is appropriate for continuously arriving data requiring low-latency handling.

Q97. CAI Connectors and CDI Integration

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

An API-driven process must call an external service and then start a CDI data pipeline without manual intervention. Which design best matches the IDMC capabilities?

A. Keep CAI and CDI isolated and require manual execution between them
B. Use CAI service connectors and integrated invocation to orchestrate external services and CDI processing
C. Use mapping transformations as a replacement for service connectors for the stated requirement
D. Use monitoring logs as the mechanism for invoking external APIs

Correct Answer: B. Use CAI service connectors and integrated invocation to orchestrate external services and CDI processing

Explanation: CAI connectors provide interaction with external services, and CAI-CDI integration supports end-to-end orchestration across application and data integration processes.

Q98. CAI Connectors and CDI Integration

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A CAI workflow needs reusable connectivity to an external application while coordinating a downstream CDI mapping. Which approach should be used?

A. Use CAI service connectors and integrated invocation to orchestrate external services and CDI processing
B. Keep CAI and CDI isolated and require manual execution between them
C. Use mapping transformations as a replacement for service connectors for the stated requirement
D. Use monitoring logs as the mechanism for invoking external APIs

Correct Answer: A. Use CAI service connectors and integrated invocation to orchestrate external services and CDI processing

Explanation: CAI connectors provide interaction with external services, and CAI-CDI integration supports end-to-end orchestration across application and data integration processes.

Q99. CAI Connectors and CDI Integration

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

The organization wants one orchestrated process across application APIs and data integration jobs. Which IDMC design pattern supports this?

A. Use CAI service connectors and integrated invocation to orchestrate external services and CDI processing
B. Keep CAI and CDI isolated and require manual execution between them
C. Use mapping transformations as a replacement for service connectors for the stated requirement
D. Use monitoring logs as the mechanism for invoking external APIs

Correct Answer: A. Use CAI service connectors and integrated invocation to orchestrate external services and CDI processing

Explanation: CAI connectors provide interaction with external services, and CAI-CDI integration supports end-to-end orchestration across application and data integration processes.

Q100. CAI Connectors and CDI Integration

Difficulty: Easy | Marks: 1 | Type: MCQ

Question Description:

A business process begins with a service request and must trigger a CDI mapping as part of the same orchestration. Which integration approach is appropriate?

A. Keep CAI and CDI isolated and require manual execution between them
B. Use mapping transformations as a replacement for service connectors for the stated requirement
C. Use monitoring logs as the mechanism for invoking external APIs
D. Use CAI service connectors and integrated invocation to orchestrate external services and CDI processing

Correct Answer: D. Use CAI service connectors and integrated invocation to orchestrate external services and CDI processing

Explanation: CAI connectors provide interaction with external services, and CAI-CDI integration supports end-to-end orchestration across application and data integration processes.

MEDIUM QUESTIONS

30 questions | 2 marks each | Scenario-based format

Q101. Data Profiling

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A retail organization is working with IDMC on a data profiling requirement. A new loyalty feed arrives from several stores with different null patterns and inconsistent category codes. The team must design cleansing rules but has no reliable data dictionary. They want evidence about completeness, distinct values, and suspicious patterns before mapping development begins. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Begin transformation design immediately using assumed formats for the stated requirement
B. Run data profiling first and use the findings to design validation and cleansing logic
C. Increase Secure Agent resources before inspecting the dataset within the integration workflow
D. Build a taskflow schedule before examining source quality as the primary design approach

Correct Answer: B. Run data profiling first and use the findings to design validation and cleansing logic

Explanation: Profiling provides factual information about the incoming data and reduces the risk of designing rules around incorrect assumptions.

Q102. Reusable Logic

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A saas organization is working with IDMC on a reusable logic requirement. Multiple client pipelines repeat the same trim, standardization, and validation steps. A rule change currently requires edits in many mappings. The team wants one maintained implementation while preserving client-specific source and target logic. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Duplicate the corrected transformations into every mapping
B. Create a separate Secure Agent for each client pipeline
C. Move the shared transformation sequence into reusable mapplet logic
D. Hardcode the common logic in mapping tasks instead

Correct Answer: C. Move the shared transformation sequence into reusable mapplet logic

Explanation: Reusable mapplet logic centralizes repeated transformation behavior and improves consistency and maintainability.

Q103. Joiner and Pushdown

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A logistics organization is working with IDMC on a joiner and pushdown requirement. Shipment and customer datasets are large and currently joined inside the integration engine. Execution time is rising because large row sets move between Snowflake and IDMC. The join is based on simple compatible keys and business logic must remain unchanged. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing

business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Use a pushdown-compatible join design so the database performs the join
B. Keep the join entirely in the mapping engine
C. Replace the join with multiple Router groups for the stated requirement
D. Increase source extraction volume before joining within the integration workflow

Correct Answer: A. Use a pushdown-compatible join design so the database performs the join

Explanation: A pushdown-compatible join leverages database processing and reduces unnecessary data movement.

Q104. Aggregator

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A telecom organization is working with IDMC on a aggregator requirement. Daily call records must be summarized to one row per subscriber and date. The existing mapping uses several expressions and duplicated branches to calculate totals. The team needs SUM, COUNT, and AVG metrics with a simpler design. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Use a Router transformation with one branch per metric
B. Use an Aggregator transformation with subscriber and date group fields
C. Use a Lookup transformation to calculate totals for the stated requirement
D. Use a Sorter transformation and load all detail rows

Correct Answer: B. Use an Aggregator transformation with subscriber and date group fields

Explanation: Aggregator is the mapping component designed for grouped calculations such as SUM, COUNT, and AVG.

Q105. Router

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A banking organization is working with IDMC on a router requirement. Transactions must follow Deposit, Withdrawal, Transfer, and Review processing paths. The current design uses repeated filters in separate branches. The team wants each row evaluated centrally and the design to remain easy to extend. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Use one Aggregator transformation per transaction type
B. Use one Router transformation with separate output-group conditions
C. Use a Joiner transformation to create branch logic
D. Use a Union transformation to test transaction categories

Correct Answer: B. Use one Router transformation with separate output-group conditions

Explanation: Router centralizes condition evaluation and supports multiple named outputs in one transformation.

Q106. Lookup Cache

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A retail organization is working with IDMC on a lookup cache requirement. Millions of sales rows require product category and brand from a small reference table. The reference data changes rarely, but the current design repeatedly accesses it during processing. The team wants to reduce repeated database calls without changing enrichment logic. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Use an appropriately cached lookup for the reference dataset
B. Use an uncached lookup for every incoming record
C. Replace the lookup with a Rank transformation
D. Re-read the reference source separately for every row

Correct Answer: A. Use an appropriately cached lookup for the reference dataset

Explanation: Caching a small, stable lookup can reduce repeated access and improve enrichment performance.

Q107. Parameterization

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A insurance organization is working with IDMC on a parameterization requirement. Development, test, and production use different schemas, file paths, and connections. The team currently maintains separate copies of the same mapping for each environment. They want one mapping design that can be promoted safely across environments. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Keep separate hardcoded mappings for every environment
B. Parameterize the environment-specific values and supply them at runtime
C. Embed production values directly in transformation expressions
D. Create one Secure Agent per parameter value

Correct Answer: B. Parameterize the environment-specific values and supply them at runtime

Explanation: Parameterization separates environment configuration from transformation logic and reduces duplicated mappings.

Q108. Taskflow Decision

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A healthcare organization is working with IDMC on a taskflow decision requirement. A pipeline validates incoming claim files before loading production tables. If validation fails, loading must stop and an alert path must run; if it succeeds, loading should continue. The current taskflow executes every step regardless of the validation result. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Run all steps sequentially and inspect results later
B. Run the load and alert paths in parallel every time
C. Add a decision step that branches on the validation outcome
D. Use a Sorter transformation to choose the taskflow path

Correct Answer: C. Add a decision step that branches on the validation outcome

Explanation: A decision step evaluates runtime results and directs orchestration to the correct branch.

Q109. Parallel Taskflow

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A telecom organization is working with IDMC on a parallel taskflow requirement. Five regional pipelines are independent but run sequentially in one taskflow. The daily processing window is becoming too long as data volume grows. A failure in one region should not delay unrelated regions. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Run the independent regional branches in parallel
B. Keep all regional tasks strictly sequential
C. Combine all regional sources with a Joiner before execution
D. Use parameterization as a substitute for parallel orchestration

Correct Answer: A. Run the independent regional branches in parallel

Explanation: Independent tasks can execute in parallel to reduce total elapsed time without violating dependencies.

Q110. Fault Handling

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A saas organization is working with IDMC on a fault handling requirement. A CAI workflow invokes an external customer API that occasionally returns temporary errors. At present, one failed call stops the entire process and requires manual restart. The team wants controlled recovery with visibility into failures. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Ignore the error and mark the process successful
B. Duplicate the entire process for every possible error
C. Configure fault handling with an appropriate retry or recovery path
D. Use a Union transformation to recover the failed API call

Correct Answer: C. Configure fault handling with an appropriate retry or recovery path

Explanation: Fault handling provides defined recovery behavior and avoids uncontrolled process termination.

Q111. CDC

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A e-commerce organization is working with IDMC on a cdc requirement. Order statuses change throughout the day, but each integration cycle reloads the complete source table. Most rows are unchanged, and processing time is increasing. The target needs timely inserts, updates, and deletes with lower source impact. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Continue full reloads for every cycle for the stated requirement

B. Use a Sorter transformation to identify all changes

C. Use result caching as the change capture mechanism

D. Use incremental change-data processing so only changed rows are handled

Correct Answer: D. Use incremental change-data processing so only changed rows are handled

Explanation: CDC/incremental processing focuses work on changed records and reduces redundant processing.

Q112. Batch vs Real-Time

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A payments organization is working with IDMC on a batch vs real-time requirement. Daily settlement files arrive once at midnight, while authorization events arrive continuously. The settlement process can tolerate hours of latency, but authorization events require near-immediate handling. The team wants ingestion modes aligned with business latency. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Use real-time ingestion for both workloads regardless of latency

B. Use batch ingestion for both workloads to simplify operations

C. Use a manual load for events and continuous flow for monthly files

D. Use batch ingestion for settlement files and real-time ingestion for authorization events

Correct Answer: D. Use batch ingestion for settlement files and real-time ingestion for authorization events

Explanation: Different latency patterns call for different ingestion modes: scheduled batch for grouped files and real-time for continuously arriving events.

Q113. Monitoring

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A telecom organization is working with IDMC on a monitoring requirement. Several CDI jobs fail intermittently at different transformations. Engineers spend too much time manually searching individual runs for root causes. They need run status, detailed execution evidence, and trend visibility. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Increase warehouse size without reviewing run evidence for the stated requirement

B. Create duplicate mappings for failed runs within the integration workflow

C. Use IDMC monitoring and execution logs to analyze failures and performance

D. Disable error logging to reduce runtime overhead as the primary design approach

Correct Answer: C. Use IDMC monitoring and execution logs to analyze failures and performance

Explanation: Monitoring and logs provide the execution evidence needed for diagnosis and trend analysis.

Q114. CAI Connector

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A travel organization is working with IDMC on a cai connector requirement. A booking workflow must call a partner pricing API and use the response in later process steps. The team currently handles the API

interaction outside CAI, which reduces visibility and reuse. They want centrally managed connectivity and error handling. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Use a Lookup transformation as the external API client
B. Use a CAI service connector within the process
C. Run the API manually outside IDMC
D. Store the API URL in a mapping expression and call it there

Correct Answer: B. Use a CAI service connector within the process

Explanation: Service connectors are designed to integrate CAI processes with external applications and APIs.

Q115. CAI-CDI

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A fintech organization is working with IDMC on a cai-cdi requirement. An API request received in CAI should validate input and then start a CDI mapping that loads Snowflake. Today an operator starts the mapping manually after the API process finishes. The business wants end-to-end automated execution. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Keep CAI and CDI isolated and retain manual handoff
B. Replace the CDI mapping with a CAI decision step
C. Invoke the CDI process from the CAI orchestration using the supported integration capability
D. Use monitoring logs to trigger the mapping manually for the stated requirement

Correct Answer: C. Invoke the CDI process from the CAI orchestration using the supported integration capability

Explanation: CAI-CDI integration supports automated orchestration across application and data integration workloads.

Q116. Hierarchical Data

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A media organization is working with IDMC on a hierarchical data requirement. Incoming JSON contains nested arrays of sessions, devices, and events. Analysts need relational fields in Snowflake while preserving parent-child relationships. Raw JSON storage alone makes downstream analysis difficult. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Discard nested elements and keep only top-level fields
B. Flatten the hierarchical structures into an analysis-friendly relational shape while preserving relationships
C. Treat each nested array as an unrelated file
D. Use a Sorter transformation to convert hierarchy into rows

Correct Answer: B. Flatten the hierarchical structures into an analysis-friendly relational shape while preserving relationships

Explanation: Hierarchical processing/flattening converts nested structures into usable relational forms while retaining relationships.

Q117. Secure Agent

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A manufacturing organization is working with IDMC on a secure agent requirement. A factory database is accessible only inside the corporate network. Cloud design assets are available, but jobs cannot directly reach the database. The team needs secure runtime connectivity without exposing the database publicly. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Expose the database publicly and avoid a runtime agent
B. Use a dashboard to connect to the on-premises database
C. Use a mapplet as the network runtime component
D. Use a Secure Agent in the appropriate runtime environment

Correct Answer: D. Use a Secure Agent in the appropriate runtime environment

Explanation: The Secure Agent executes jobs and provides connectivity between IDMC and protected systems.

Q118. Runtime Environment

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A banking organization is working with IDMC on a runtime environment requirement. The organization has multiple Secure Agents assigned to different workload groups. A mapping task must run only on the agents approved for finance integrations. The team needs a supported way to select those runtime resources. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Select a project folder as the runtime resource
B. Assign the mapping task to the runtime environment containing the approved agents
C. Store the agent names in an Expression transformation
D. Use a connection object as the execution environment

Correct Answer: B. Assign the mapping task to the runtime environment containing the approved agents

Explanation: Runtime environments group Secure Agent resources for job execution.

Q119. Connection

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A education organization is working with IDMC on a connection requirement. Several mappings must read the same cloud database. Developers currently repeat endpoint details in separate places, creating maintenance risk. The team wants centrally managed access settings that can be reused. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Create and reuse a connection containing endpoint and authentication properties

B. Hardcode endpoint details in every mapping expression for the stated requirement

C. Store credentials in taskflow decision conditions within the integration workflow

D. Create a separate mapplet for database authentication as the primary design approach

Correct Answer: A. Create and reuse a connection containing endpoint and authentication properties

Explanation: Connections centralize the properties IDMC uses to access a source or target system.

Q120. Mapping vs Task

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A healthcare organization is working with IDMC on a mapping vs task requirement. A mapping contains validated source-to-target logic used by several scheduled jobs. Each job needs different runtime parameters and schedules, but the transformation design should remain the same. The team wants clear separation of design and execution configuration. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Copy the mapping for every schedule for the stated requirement

B. Reuse the mapping and create mapping task configurations for the different executions

C. Put all transformation logic only in the task object

D. Use a connection object instead of a mapping task

Correct Answer: B. Reuse the mapping and create mapping task configurations for the different executions

Explanation: Mappings define data logic, while mapping tasks provide runnable configuration.

Q121. Pushdown Filter

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A retail organization is working with IDMC on a pushdown filter requirement. A source contains hundreds of columns and years of history, but downstream reporting uses only recent rows and a subset of fields. The mapping currently reads far more data than necessary before filtering. The team wants to reduce data movement with minimal redesign. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Read the full dataset and filter only before the target

B. Apply compatible filtering/projection as close to the source database as possible through pushdown

C. Increase Secure Agent memory and keep the same flow

D. Disable database-side processing to simplify execution for the stated requirement

Correct Answer: B. Apply compatible filtering/projection as close to the source database as possible through pushdown

Explanation: Early database-side filtering reduces transferred data and can improve overall execution efficiency.

Q122. Sorter

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A logistics organization is working with IDMC on a sorter requirement. A legacy export must be delivered in customer and shipment-date order. The target consumer depends on deterministic row ordering. No grouping or filtering requirement exists. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Use an Aggregator transformation to impose order
B. Use a Lookup transformation to control sequence
C. Use a Router transformation to sort each category
D. Use a Sorter transformation with the required sort keys

Correct Answer: D. Use a Sorter transformation with the required sort keys

Explanation: Sorter is the component intended to order rows by configured keys.

Q123. Union

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A insurance organization is working with IDMC on a union requirement. Two claim feeds have the same target structure and need to be processed as one stream. There is no business-key match requirement between the feeds. The team wants to append the rows before common downstream transformations. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Use a Union transformation to combine the compatible row streams
B. Use a Joiner transformation with a fabricated key
C. Use a Lookup to append rows from the second feed
D. Use a Rank transformation to merge the pipelines

Correct Answer: A. Use a Union transformation to combine the compatible row streams

Explanation: Union appends compatible row sets vertically without key-based matching.

Q124. Rank

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A retail organization is working with IDMC on a rank requirement. A dashboard feed needs the top ten products by revenue in each region. Sorting alone would still pass every product downstream. The team needs only a limited set of highest-ranked rows. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Use only a Sorter transformation and keep all rows
B. Use a Router transformation with one output per product
C. Use a Lookup transformation to determine row rank
D. Use a Rank transformation configured for the required top records

Correct Answer: D. Use a Rank transformation configured for the required top records

Explanation: Rank selects top or bottom records based on configured ranking criteria.

Q125. Performance Tuning

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A saas organization is working with IDMC on a performance tuning requirement. A mapping has gradually slowed as data volumes increased. The team sees heavy engine processing, repeated reference access, and more data movement than expected. They want evidence-based optimization rather than simply adding resources. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Review logs and execution design, then optimize pushdown, lookup strategy, and unnecessary data movement

B. Increase all compute sizes before reviewing execution behavior for the stated requirement

C. Disable monitoring so runtime overhead is lower within the integration workflow

D. Duplicate the mapping into smaller copies without analysis as the primary design approach

Correct Answer: A. Review logs and execution design, then optimize pushdown, lookup strategy, and unnecessary data movement

Explanation: Performance tuning should target the actual bottleneck using execution evidence and appropriate design changes.

Q126. Taskflow Command

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A manufacturing organization is working with IDMC on a taskflow command requirement. A pipeline must run a supported system utility after data export completes. The action is outside the data mapping itself but belongs to the orchestration. The team wants the taskflow to control the operation and its status. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Put the shell command inside an Aggregator transformation

B. Use a Joiner transformation to invoke the utility

C. Use a Lookup transformation as the command executor

D. Use a taskflow command step for the system-level action

Correct Answer: D. Use a taskflow command step for the system-level action

Explanation: Command steps are the appropriate orchestration mechanism for supported system-level commands.

Q127. Event-Driven Taskflow

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A travel organization is working with IDMC on a event-driven taskflow requirement. A process should start when a defined application event arrives rather than on a fixed daily schedule. The event may arrive at unpredictable times. The team wants native event-driven orchestration. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Use the relevant taskflow trigger or listener capability for the event

B. Poll with a monthly batch schedule for the stated requirement

C. Use a Sorter transformation to wait for the event

D. Create a mapping task with no trigger and run it manually

Correct Answer: A. Use the relevant taskflow trigger or listener capability for the event

Explanation: Triggers/listeners support event-driven taskflow execution.

Q128. Error Quarantine

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A financial services organization is working with IDMC on a error quarantine requirement. A transaction pipeline contains some invalid records with missing mandatory fields. The business wants valid rows to continue while failed rows are retained for investigation. Stopping the whole pipeline is unacceptable. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Load invalid rows into the production target without validation for the stated requirement
B. Route invalid rows to a separate error or quarantine flow while valid rows continue
C. Stop the entire pipeline on the first invalid record within the integration workflow
D. Discard invalid rows without any trace as the primary design approach

Correct Answer: B. Route invalid rows to a separate error or quarantine flow while valid rows continue

Explanation: Separating invalid records preserves throughput while maintaining traceability for remediation.

Q129. CAI Parallelism

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A healthcare organization is working with IDMC on a cai parallelism requirement. A CAI process calls several independent external services before a final consolidation step. The calls currently run one after another and dominate response time. Only the consolidation step depends on all results. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Keep all service calls sequential regardless of dependency for the stated requirement
B. Cache old responses instead of executing current calls within the integration workflow
C. Run the independent service calls in parallel and join the flow before consolidation
D. Remove the consolidation dependency entirely as the primary design approach

Correct Answer: C. Run the independent service calls in parallel and join the flow before consolidation

Explanation: Parallel execution is appropriate for independent steps and can reduce end-to-end latency while preserving the final dependency.

Q130. Expression Transformation

Difficulty: Medium | Marks: 2 | Type: MCQ

Question Description:

A healthcare organization is working with IDMC on a expression transformation requirement. A provider feed contains raw duration, quantity, and rate fields that must produce several derived business columns. The current mapping uses multiple duplicated calculations across downstream branches. The team wants one clear place for row-level derivation logic before later filtering and loading. The team wants a maintainable solution that uses the platform capabilities appropriately. They want to avoid unnecessary

manual work and preserve the existing business intent. The architect is reviewing the design before implementation. Based on this scenario, which decision is most appropriate?

A. Use an Expression transformation to calculate and standardize the derived fields
B. Use an Aggregator transformation for every row-level formula
C. Use a Router transformation to calculate numerical expressions
D. Use a Union transformation to derive new columns

Correct Answer: A. Use an Expression transformation to calculate and standardize the derived fields

Explanation: Expression is designed for row-level calculations, conversions, and derived output fields without grouping records.

HARD QUESTIONS

30 questions | 2 marks each | Scenario-based format

Q131. Pushdown Compatibility

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A retail organization operates a production IDMC solution involving pushdown compatibility. A high-volume order pipeline uses filters, joins, expressions, and aggregations before loading Snowflake. Most operations still execute in the integration engine even though Snowflake has available compute. The profile shows heavy data movement and long engine-side processing. Some expressions use functions that are not compatible with database-side execution, and business logic cannot change materially. The architect wants maximum pushdown while preserving correctness and maintainability. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Force full pushdown without checking function compatibility for the stated requirement

B. Refactor only the incompatible expressions into pushdown-compatible logic so the eligible operations can execute in Snowflake

C. Disable pushdown and scale the Secure Agent instead within the integration workflow

D. Stage every transformation result in separate files before loading as the primary design approach

Correct Answer: B. Refactor only the incompatible expressions into pushdown-compatible logic so the eligible operations can execute in Snowflake

Explanation: Pushdown works best when transformation logic is compatible with the target database; refactoring incompatibilities can move more work to Snowflake without changing business intent.

Q132. CDC Recovery

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A banking organization operates a production IDMC solution involving cdc recovery. A transaction feed uses incremental change processing into Snowflake. One run fails after some downstream work completes but before the watermark is safely advanced. Restarting from the wrong point could either duplicate changes or skip transactions. The source continues receiving new updates while recovery is performed. The team needs a restartable design with accurate incremental boundaries and traceability. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Advance the checkpoint before processing begins for the stated requirement within the integration workflow

B. Persist and advance the incremental checkpoint only after successful processing, then restart from the last committed checkpoint

C. Use the current clock time as the restart point regardless of prior success

D. Switch to full reloads after every failure without preserving incremental state as the primary design approach

Correct Answer: B. Persist and advance the incremental checkpoint only after successful processing, then restart from the last committed checkpoint

Explanation: A committed checkpoint tied to successful processing prevents gaps and makes CDC restart behavior deterministic.

Q133. Parallel Dependency Control

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A telecom organization operates a production IDMC solution involving parallel dependency control. A taskflow has eight regional mappings, two shared validation tasks, and one final consolidation. Regional mappings are independent, but consolidation requires every successful regional output. The current design is fully sequential and exceeds the batch window. Running everything in parallel would violate the final dependency and complicate failure handling. The architect wants shorter elapsed time without losing dependency control. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Keep all tasks sequential to preserve simplicity for the stated requirement

B. Run consolidation in parallel with regional mappings within the integration workflow

C. Run independent regional mappings in parallel, then gate consolidation on completion of the required branches

D. Remove validation dependencies and rely on target checks later as the primary design approach

Correct Answer: C. Run independent regional mappings in parallel, then gate consolidation on completion of the required branches

Explanation: The correct design combines parallelism for independent work with explicit dependency control before the consolidation step.

Q134. Lookup Strategy

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A e-commerce organization operates a production IDMC solution involving lookup strategy. A product lookup contains 200,000 rows and changes once per day. The fact stream contains tens of millions of rows and repeatedly requests the same product keys. Direct database lookups create excessive calls, but keeping stale data after the daily refresh is unacceptable. The team needs fast enrichment and predictable refresh behavior. The architect is choosing a lookup strategy for the daily batch. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Use an uncached database lookup for every fact row for the stated requirement

B. Use a persistent cache that is never refreshed within the integration workflow

C. Replace the lookup with a Router transformation based on product keys as the primary design approach

D. Use a lookup cache refreshed for each batch so repeated keys are resolved locally using current reference data

Correct Answer: D. Use a lookup cache refreshed for each batch so repeated keys are resolved locally using current reference data

Explanation: A batch-scoped cache reduces repeated database access while still allowing the reference data to be refreshed each run.

Q135. Reusable Governance

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A saas organization operates a production IDMC solution involving reusable governance. Dozens of mappings use shared validation logic owned by a central data-quality team. Business units also need limited flexibility for client-specific rules. Copying the shared rules has caused drift, while one giant shared component would make every change risky. The team wants governed reuse with clear ownership and extension points. The architect is reviewing modular design. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Copy the entire validation flow into every client mapping for the stated requirement
B. Move all client-specific rules into one global mapplet within the integration workflow
C. Hardcode shared rules in mapping task properties as the primary design approach
D. Place stable common rules in reusable mapplet logic and keep client-specific logic outside the shared component

Correct Answer: D. Place stable common rules in reusable mapplet logic and keep client-specific logic outside the shared component

Explanation: Separating stable shared logic from client-specific logic provides reuse without making the reusable component overly coupled.

Q136. Secure Agent Resilience

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A manufacturing organization operates a production IDMC solution involving secure agent resilience. Two factories depend on on-premises data sources and run critical nightly integrations. A single Secure Agent currently handles all jobs, creating a runtime single point of failure. The business requires better runtime availability without exposing source systems publicly. Connections and mappings should not need redesign when an agent host fails. The platform team is reviewing runtime architecture. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Use a runtime environment with multiple appropriately configured Secure Agents so jobs have resilient execution resources
B. Expose the databases publicly and remove the agents for the stated requirement
C. Create duplicate mappings tied to different hardcoded hosts within the integration workflow
D. Use a taskflow decision step as a substitute for runtime redundancy

Correct Answer: A. Use a runtime environment with multiple appropriately configured Secure Agents so jobs have resilient execution resources

Explanation: Grouping multiple Secure Agents in the runtime environment improves execution resilience while preserving the same integration design.

Q137. Environment Promotion

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A insurance organization operates a production IDMC solution involving environment promotion. The same integration solution moves from development to test and production. Connections, schemas, file locations, and credentials differ in each environment. Current promotions require manual edits to mappings, which has caused production mistakes. The team wants controlled promotion with minimal design changes and repeatable configuration. Security-sensitive values must not be hardcoded. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Maintain separate independently edited mappings for each environment

B. Hardcode production values before promoting the development asset

C. Store all environment differences in transformation expressions for the stated requirement

D. Parameterize environment-specific properties and supply approved runtime configuration during deployment or execution

Correct Answer: D. Parameterize environment-specific properties and supply approved runtime configuration during deployment or execution

Explanation: Parameterization separates environment configuration from mapping design and supports safer promotion.

Q138. CAI Fault Isolation

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A fintech organization operates a production IDMC solution involving cai fault isolation. A CAI payment workflow calls identity, risk, notification, and ledger services. A notification outage should not roll back a successful ledger action, while identity or risk failure must stop the payment path. The current workflow treats every service failure identically. The business needs different recovery behavior based on service criticality. Monitoring must show which branch failed and what recovery occurred. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Apply one ignore-error rule to every service call for the stated requirement within the integration workflow

B. Stop the entire workflow for any noncritical notification failure as the primary design approach

C. Define fault handling per service path, with stop/recovery for critical calls and isolated retry or alternate handling for noncritical calls

D. Remove all fault handling and rely on manual support for this processing scenario in the current pipeline design

Correct Answer: C. Define fault handling per service path, with stop/recovery for critical calls and isolated retry or alternate handling for noncritical calls

Explanation: Granular fault handling lets the process preserve critical correctness while isolating failures that should not invalidate completed work.

Q139. Hierarchical Flattening

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A healthcare organization operates a production IDMC solution involving hierarchical flattening. A JSON claim document contains claim, member, provider, and repeating procedure arrays. Analysts need

relational outputs while maintaining the association between each procedure and its parent claim. A naive flatten creates duplicate or orphaned relationships. The pipeline must remain scalable as nested arrays grow. The team wants a structured design rather than storing raw JSON only. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Flatten the hierarchy with preserved parent identifiers so repeated child records remain linked to the correct parent

B. Flatten all arrays without carrying parent identifiers for the stated requirement within the integration workflow

C. Discard the repeating procedure arrays during transformation as the primary design approach

D. Store each nested object independently without relationship keys for this processing scenario

Correct Answer: A. Flatten the hierarchy with preserved parent identifiers so repeated child records remain linked to the correct parent

Explanation: Hierarchical processing must preserve parent-child keys during flattening to maintain data integrity.

Q140. Monitoring and Root Cause

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A telecom organization operates a production IDMC solution involving monitoring and root cause. A mapping shows unpredictable runtime spikes on some days. Source volume varies moderately, but spikes do not always correlate with row count. Execution logs show occasional slow lookups and periods of heavy engine processing. The team has been increasing compute reactively without proving the bottleneck. They need a repeatable diagnostic approach before tuning. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Increase all compute resources permanently for the stated requirement within the integration workflow

B. Disable monitoring to remove any logging overhead as the primary design approach

C. Correlate monitoring logs, step-level execution behavior, lookup activity, and data volume before selecting a targeted optimization

D. Rewrite every mapping transformation without analyzing evidence for this processing scenario

Correct Answer: C. Correlate monitoring logs, step-level execution behavior, lookup activity, and data volume before selecting a targeted optimization

Explanation: Evidence-based tuning identifies the actual bottleneck and avoids expensive changes that may not address the problem.

Q141. Pushdown Join Design

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A logistics organization operates a production IDMC solution involving pushdown join design. A join between shipments and routes is eligible for database execution except for a custom expression embedded in the join condition. Because of that expression, the join stays in the IDMC engine and moves large datasets. The same custom expression can be derived earlier using database-compatible functions.

The team wants to preserve the business rule while enabling more pushdown. The architect is deciding where to place the derivation. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Derive the compatible join key before the join and use a simple database-compatible join condition
B. Keep the complex expression inside the join and force pushdown
C. Move the join entirely to the mapping engine for the stated requirement
D. Replace the join with a Router and Union combination within the integration workflow

Correct Answer: A. Derive the compatible join key before the join and use a simple database-compatible join condition

Explanation: Simplifying the join predicate with pre-derived compatible keys improves pushdown eligibility without changing the matching rule.

Q142. Error Quarantine and Restart

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A financial services organization operates a production IDMC solution involving error quarantine and restart. A daily load processes millions of records, with a small number failing validation. The pipeline currently stops on the first invalid row and must be rerun from the beginning. Valid data should not be delayed, but failed records must be traceable and reprocessable after correction. The team also wants reruns to avoid duplicating already accepted rows. The architect needs a resilient operating model. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Stop the entire run for any invalid record for the stated requirement
B. Discard rejected rows and rerun the complete source within the integration workflow
C. Load invalid rows into production and flag them later as the primary design approach
D. Route failed rows to a quarantine path, commit valid processing, and design controlled reprocessing of corrected rejects

Correct Answer: D. Route failed rows to a quarantine path, commit valid processing, and design controlled reprocessing of corrected rejects

Explanation: A quarantine/reprocessing pattern preserves throughput, traceability, and controlled recovery without repeating valid work.

Q143. CAI-CDI Transaction Boundary

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A retail organization operates a production IDMC solution involving cai-cdi transaction boundary. An API request in CAI starts a CDI pipeline that writes business data. The API caller needs a meaningful completion status, but the data job may run longer than the initial request window. The current design blocks the API until the data pipeline finishes and causes timeouts. The business can accept asynchronous processing if status is trackable. The solution should remain orchestrated and observable. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Use CAI to trigger the CDI work asynchronously and return a trackable process status rather than blocking for full completion

B. Keep the synchronous API call open until every data step finishes for the stated requirement

C. Remove CDI orchestration and require manual job starts within the integration workflow as the primary design approach

D. Return success immediately without recording any downstream status for this processing scenario in the current pipeline design

Correct Answer: A. Use CAI to trigger the CDI work asynchronously and return a trackable process status rather than blocking for full completion

Explanation: Asynchronous orchestration avoids request timeout while retaining traceable process execution.

Q144. Batch Window Optimization

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A insurance organization operates a production IDMC solution involving batch window optimization. A nightly process contains ingestion, reference enrichment, several independent regional transforms, and one final aggregate. The batch misses its SLA by forty minutes. Profiling shows ingestion and enrichment are shared prerequisites, regional transforms are independent, and the aggregate requires all regional outputs. Simply scaling every component would increase cost without addressing the serial design. The team needs a structural optimization. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Run the aggregate at the same time as regional transforms for the stated requirement

B. Keep the entire workflow sequential and only increase compute within the integration workflow

C. Keep shared prerequisites sequential, execute independent regional transforms in parallel, then run the final aggregate after all branches complete

D. Duplicate ingestion in every regional branch as the primary design approach for this processing scenario

Correct Answer: C. Keep shared prerequisites sequential, execute independent regional transforms in parallel, then run the final aggregate after all branches complete

Explanation: This hybrid taskflow preserves dependencies while exploiting safe parallelism to reduce elapsed time.

Q145. Dynamic Runtime Values

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A media organization operates a production IDMC solution involving dynamic runtime values. A reusable mapping loads different content tables depending on tenant and date. The current design constructs separate mappings for each tenant, and every new tenant requires code duplication. Table names and file paths vary, but transformation logic is stable. The business expects frequent onboarding of new tenants. The team wants controlled runtime flexibility. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Clone the mapping for each tenant and hardcode values

B. Use parameters for tenant-specific object names and paths while maintaining one transformation design

C. Store tenant names in a Sorter transformation for the stated requirement

D. Create a new Secure Agent for each tenant within the integration workflow

Correct Answer: B. Use parameters for tenant-specific object names and paths while maintaining one transformation design

Explanation: Parameterization is intended for runtime variation of configuration values without duplicating core mapping logic.

Q146. Real-Time vs Batch Boundary

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A payments organization operates a production IDMC solution involving real-time vs batch boundary. The platform receives high-volume authorization events continuously and reconciliation files every night. Authorization decisions need low latency, while reconciliation can run in a scheduled window. The team is considering one ingestion pattern for operational simplicity. Using the wrong pattern could either increase latency or operational overhead. The architect wants each workload aligned to its service-level requirement. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Use real-time ingestion for authorization events and scheduled batch ingestion for reconciliation files

B. Use batch for authorizations and real-time for nightly reconciliation

C. Use real-time processing for every source regardless of latency needs

D. Use manual ingestion for both workloads for the stated requirement

Correct Answer: A. Use real-time ingestion for authorization events and scheduled batch ingestion for reconciliation files

Explanation: The source supports distinct batch and real-time patterns; matching them to latency requirements gives the appropriate design.

Q147. Data Profiling Governance

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A healthcare organization operates a production IDMC solution involving data profiling governance. Several source systems use the same field name but with different value conventions. Teams have been writing cleansing rules independently, producing inconsistent downstream interpretation. A new integration program wants standardized rules based on observed source behavior. The architects need objective evidence before defining common transformations. The solution should also make future source onboarding easier. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Define one global rule without examining any source data for the stated requirement

B. Skip profiling and rely only on field names within the integration workflow

C. Create separate unrelated cleansing logic for every source as the primary design approach

D. Profile representative sources first, document the observed patterns, and base shared validation rules on those findings

Correct Answer: D. Profile representative sources first, document the observed patterns, and base shared validation rules on those findings

Explanation: Profiling supplies the evidence needed to standardize transformation rules across heterogeneous sources.

Q148. Connection Security

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A banking organization operates a production IDMC solution involving connection security. Many mappings access the same protected database. Developers have started placing endpoint and login details in mapping notes and local files. The security team wants centralized, reusable connection configuration with controlled access. The integration logic should not contain raw credentials. Operational changes to the endpoint should require minimal mapping updates. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Embed credentials directly in every mapping expression for the stated requirement
B. Store the credentials in taskflow branch labels within the integration workflow
C. Create duplicate mappings whenever the endpoint changes as the primary design approach
D. Use centrally managed IDMC connection objects and reference them from mappings/tasks

Correct Answer: D. Use centrally managed IDMC connection objects and reference them from mappings/tasks

Explanation: Connection objects centralize endpoint/authentication properties and reduce credential sprawl in integration logic.

Q149. Taskflow Conditional Recovery

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A education organization operates a production IDMC solution involving taskflow conditional recovery. A student-load workflow validates files, loads good records, and sends status notifications. If validation finds only warnings, loading should continue with a warning notification. If validation finds critical errors, loading must stop and a remediation branch must run. The current flow has only pass/fail behavior. The team wants native orchestration without custom external scripts. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Use a single sequential path for all outcomes
B. Use decision logic with distinct branches for pass, warning, and critical outcomes
C. Run every branch in parallel regardless of validation result
D. Use a Joiner transformation to select the taskflow path

Correct Answer: B. Use decision logic with distinct branches for pass, warning, and critical outcomes

Explanation: Decision steps allow taskflows to branch based on runtime conditions and support multiple outcomes.

Q150. Reference Enrichment Scalability

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A retail organization operates a production IDMC solution involving reference enrichment scalability. A product-reference dataset grows from thousands to millions of rows and is updated more frequently. The previous fully cached lookup now consumes excessive memory and refresh time. Most fact rows still need enrichment by product key. The team wants to reassess the lookup design rather than keep the old configuration blindly. Database-side processing is available. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Keep the large lookup fully cached regardless of growth for the stated requirement

B. Use an uncached row-by-row lookup without analysis within the integration workflow

C. Replace the reference enrichment with a Rank transformation as the primary design approach

D. Evaluate pushdown or database-side join/enrichment instead of assuming the existing full-cache lookup remains optimal

Correct Answer: D. Evaluate pushdown or database-side join/enrichment instead of assuming the existing full-cache lookup remains optimal

Explanation: As reference size and update frequency change, the previous cache strategy may no longer be optimal; database-side enrichment can reduce agent pressure and movement.

Q151. Reusable Component Versioning

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A saas organization operates a production IDMC solution involving reusable component versioning. A reusable cleansing component is used by twenty mappings. A proposed change is valid for new clients but could alter historical behavior for existing ones. Updating the shared object in place would affect all consumers at once. The team needs controlled adoption and regression safety. Governance requires predictable change impact. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Modify the shared component in place without testing consumers for the stated requirement

B. Introduce a managed new version of the reusable logic and migrate consumers deliberately after validation

C. Duplicate random copies inside each mapping within the integration workflow

D. Remove reuse and hardcode all rules locally as the primary design approach

Correct Answer: B. Introduce a managed new version of the reusable logic and migrate consumers deliberately after validation

Explanation: Controlled versioning of reusable logic reduces unintended cross-mapping impact while retaining maintainability.

Q152. Orchestration Observability

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A travel organization operates a production IDMC solution involving orchestration observability. A taskflow coordinates five tasks and two external calls. Users report occasional delays, but the final status alone

does not show which step is responsible. The team needs actionable evidence for duration, failures, and branch behavior. They also want to distinguish slow external calls from slow data tasks. The design should use platform monitoring rather than custom spreadsheets. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Only record the final success/failure status for the stated requirement
B. Increase all task timeouts without measuring step durations within the integration workflow
C. Disable logs to reduce noise as the primary design approach
D. Use detailed taskflow/run monitoring and step-level logs to isolate delay and failure locations

Correct Answer: D. Use detailed taskflow/run monitoring and step-level logs to isolate delay and failure locations

Explanation: Step-level monitoring provides the observability needed to diagnose orchestration delays and failures.

Q153. CAI Connector Reuse

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A fintech organization operates a production IDMC solution involving cai connector reuse. Several CAI processes call the same customer-verification API with the same authentication and endpoint. Each process currently implements the call differently, creating inconsistent error handling. The API team may change the endpoint in the future. The integration team wants reusable service access with centralized maintenance. The process-specific business logic should remain separate. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Use a reusable service connector/configuration for the API and keep workflow-specific logic in each process
B. Hardcode the endpoint separately in every process for the stated requirement
C. Use CDI mapplets as the API communication mechanism within the integration workflow
D. Copy the entire CAI process for each use case as the primary design approach

Correct Answer: A. Use a reusable service connector/configuration for the API and keep workflow-specific logic in each process

Explanation: Reusable connector configuration centralizes service access while allowing each process to retain its own orchestration logic.

Q154. Source Filtering vs Engine Filtering

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A retail organization operates a production IDMC solution involving source filtering vs engine filtering. A source table holds five years of sales, but a daily pipeline only needs yesterday's rows. The current mapping extracts all years and applies a filter later in the engine. Network transfer and runtime continue to rise as history grows. The business rule is a simple date predicate supported by the source database. The team wants the smallest design change with the largest reduction in movement. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing

several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Keep extracting all history and increase Secure Agent memory
B. Filter only after all joins and aggregations complete for the stated requirement
C. Push the date predicate to the source/database so only required rows are read
D. Disable pushdown so execution behavior is uniform within the integration workflow

Correct Answer: C. Push the date predicate to the source/database so only required rows are read

Explanation: Applying a supported predicate at the source reduces extracted volume and downstream processing.

Q155. Snowflake Concurrency Context

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A analytics organization operates a production IDMC solution involving snowflake concurrency context. Multiple IDMC loads and dashboard queries share Snowflake during peak hours. The IDMC mappings themselves are efficient, but analytics queries queue heavily when concurrency spikes. Increasing warehouse size helps individual query speed but does not fully address simultaneous demand. The architecture needs a Snowflake-side concurrency decision without changing IDMC mappings. The source material uses multi-cluster warehouses for this pattern. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Use result caching as the only concurrency strategy
B. Force all workloads into one sequential schedule
C. Use multi-cluster warehouse behavior for the concurrency-heavy Snowflake workload
D. Disable warehouse scaling and increase IDMC engine processing

Correct Answer: C. Use multi-cluster warehouse behavior for the concurrency-heavy Snowflake workload

Explanation: Multi-cluster behavior addresses concurrent query demand, which is distinct from making an individual query faster.

Q156. Snowflake Idle Cost Context

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A saas organization operates a production IDMC solution involving snowflake idle cost context. Several Snowflake warehouses support scheduled IDMC jobs but remain idle for long periods between runs. Compute cost continues during inactivity because automatic suspension is not configured. The team wants jobs to resume when needed without manual operations. No mapping redesign is required. The source material explicitly uses auto-suspend for this cost pattern. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Configure automatic warehouse suspension during inactivity with resume behavior as needed
B. Keep warehouses continuously running for faster startup for the stated requirement
C. Change every mapping to use a different connection

D. Use clustering keys to control compute billing within the integration workflow

Correct Answer: A. Configure automatic warehouse suspension during inactivity with resume behavior as needed

Explanation: Auto-suspend is the source-supported mechanism for reducing idle warehouse compute cost.

Q157. Snowflake Query Cache Context

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A media organization operates a production IDMC solution involving snowflake query cache context. A dashboard repeatedly runs identical Snowflake queries against data that changes only once per day. Some users see very fast repeated responses, while responses slow after the daily load changes underlying data. The team wants to understand the condition rather than force caching blindly. Warehouse size is unchanged. The source material distinguishes valid result reuse from changed-data behavior. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Assume cached results remain valid after source data changes for the stated requirement

B. Increase warehouse size to force reuse of old results within the integration workflow

C. Rely on result-cache reuse only while the query is eligible and underlying data remains unchanged

D. Use clustering to keep stale query results available as the primary design approach

Correct Answer: C. Rely on result-cache reuse only while the query is eligible and underlying data remains unchanged

Explanation: Result-cache reuse depends on cache eligibility and unchanged underlying data; changes invalidate the previous result.

Q158. Snowflake Pruning Context

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A fintech organization operates a production IDMC solution involving snowflake pruning context. A very large Snowflake transaction table is frequently filtered by date and region. Queries are scanning more micro-partitions than expected after substantial data evolution. The IDMC pipeline is not the main bottleneck; the issue is data access efficiency in Snowflake. The team wants to improve pruning with minimal query changes. The source material discusses clustering and pruning for this pattern. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Move all filtering back into the IDMC engine for the stated requirement

B. Use a Router transformation to reorganize Snowflake storage within the integration workflow

C. Use an appropriate clustering strategy on frequently filtered columns when the table benefits from it

D. Increase taskflow parallelism to reduce partition scans as the primary design approach

Correct Answer: C. Use an appropriate clustering strategy on frequently filtered columns when the table benefits from it

Explanation: Clustering on appropriate filter columns can improve micro-partition pruning for large Snowflake tables.

Q159. Full Pipeline Recovery

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A logistics organization operates a production IDMC solution involving full pipeline recovery. A multi-step workflow ingests files, transforms data, and loads Snowflake. An upstream ingestion task succeeds, a middle transform fails, and downstream tasks must not run. The input file should not be ingested twice during recovery, and operators need a clear restart point. The workflow currently has no explicit recovery design. The architect needs controlled fault handling and restartability. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Rerun every step blindly from the beginning for the stated requirement

B. Ignore the failed transform and continue downstream within the integration workflow

C. Use taskflow fault handling plus idempotent/checkpointed steps so recovery resumes from a safe known point

D. Delete all monitoring history before retrying as the primary design approach

Correct Answer: C. Use taskflow fault handling plus idempotent/checkpointed steps so recovery resumes from a safe known point

Explanation: Fault handling combined with restart-safe task design prevents duplicate work and provides a clear recovery path.

Q160. Pushdown Tradeoff

Difficulty: Hard | Marks: 2 | Type: MCQ

Question Description:

A healthcare organization operates a production IDMC solution involving pushdown tradeoff. A mapping can push most logic to Snowflake, but one business rule depends on functionality unavailable for pushdown. Forcing database execution would change the rule, while keeping all logic in the engine creates unnecessary data movement. The team wants the best compromise. The design must preserve exact business semantics. The architect is choosing between all-or-nothing execution approaches. The workload is business-critical and the team must avoid introducing a workaround that only hides the underlying design issue. Operational traceability, maintainability, and predictable recovery are required. The architect is comparing several alternatives and wants the design that best satisfies all stated conditions. Based on the scenario, which decision is most appropriate?

A. Change the business rule so full pushdown becomes possible for the stated requirement

B. Disable all pushdown because one operation is incompatible within the integration workflow

C. Push compatible operations to Snowflake and keep only the incompatible rule in the integration engine

D. Force the incompatible operation to the database even if semantics change

Correct Answer: C. Push compatible operations to Snowflake and keep only the incompatible rule in the integration engine

Explanation: Selective/partial pushdown preserves correctness while still moving eligible work closer to the data.